

ALY6040_Week1_EDA_Group6

July 20, 2022

1 H&M Recommender System - Exploratory Data Analysis

1.1 ALY6040 - Data Mining

1.2 Group 6 - Harkirat Singh, Harshit Gaur, Puneet Madan, Akash Raj

1.3 Introduction

The H&M Group is a collection of brands and companies with about 4,850 physical locations and 53 online marketplaces. Customers can browse a wide assortment of products in our online store. However, if there are too many options, clients could not find what they are looking for or what intrigues them right away, which could prevent them from making a purchase. Product recommendations are essential for improving the buying experience. More importantly, assisting consumers in making sound decisions benefits sustainability since it lowers returns and, as a result, lowers transportation-related emissions. We are going to develop product recommendations based on data from previous transactions, as well as from customer and product meta data.

1.4 About the dataset

The dataset contains the following files.

- images/ - a folder of images corresponding to each article_id. The images are placed in subfolders starting with the first three digits of the article_id.
- articles.csv - CSV file containing the detailed metadata for each article_id available for purchase.
- customers.csv - CSV file containing the metadata for each customer_id in dataset.
- transactions_train.csv - CSV file containing the transactions data. It consists of the purchases of each customer for each date, as well as additional information. Duplicate rows correspond to multiple purchases of the same item.

We are going to start off our analysis by exploring the data, understanding the meaning and significance of the attributes involved properly, and perform necessary actions to streamline our process for the road to our goal.

1.4.1 Importing data from kaggle

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import json
```

```
import matplotlib.pyplot as plt
pd.set_option('display.max_columns', 40)
```

```
[ ]: # Label Encoding and One-Hot Encoding Libraries
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder

# Standardization and Normalization Libraries
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler

# SciPy.Stats for Plotting
import scipy.stats as stats
import pylab
```

```
[ ]: !mkdir ~/.kaggle
```

```
[ ]: json_string = {"username":"mrnerd","key":"e90faefa3b2a5f6183c87004e6f7dd56"}
with open('/root/.kaggle/kaggle.json', 'w', encoding='utf-8') as f:
    json.dump(json_string, f, ensure_ascii=False, indent=4)
```

```
[ ]: !chmod 600 /root/.kaggle/kaggle.json
```

1.4.2 Downloading and Unzipping the whole dataset (if not using images, then use another script after this)

```
[ ]: #! kaggle competitions download -c h-and-m-personalized-fashion-recommendations
```

Downloading h-and-m-personalized-fashion-recommendations.zip to /content

100% 28.7G/28.7G [03:56<00:00, 149MB/s]

100% 28.7G/28.7G [03:56<00:00, 130MB/s]

```
[ ]: #! unzip /content/h-and-m-personalized-fashion-recommendations.zip
```

1.4.3 Downloading and Unzipping specific files from Kaggle Dataset

```
[ ]: ! kaggle competitions download -c h-and-m-personalized-fashion-recommendations
    ↪ -f articles.csv
! kaggle competitions download -c h-and-m-personalized-fashion-recommendations
    ↪ -f customers.csv
! kaggle competitions download -c h-and-m-personalized-fashion-recommendations
    ↪ -f sample_submission.csv
! kaggle competitions download -c h-and-m-personalized-fashion-recommendations
    ↪ -f transactions_train.csv
```

Downloading articles.csv.zip to /content

0% 0.00/4.26M [00:00<?, ?B/s]

100% 4.26M/4.26M [00:00<00:00, 97.2MB/s]

```

Downloading customers.csv.zip to /content
 79% 77.0M/97.9M [00:00<00:00, 90.9MB/s]
100% 97.9M/97.9M [00:00<00:00, 117MB/s]
Downloading sample_submission.csv.zip to /content
 87% 44.0M/50.3M [00:00<00:00, 169MB/s]
100% 50.3M/50.3M [00:00<00:00, 154MB/s]
Downloading transactions_train.csv.zip to /content
 96% 562M/584M [00:05<00:00, 98.0MB/s]
100% 584M/584M [00:05<00:00, 118MB/s]

```

```
[ ]: ! unzip /content/articles.csv.zip
! unzip /content/customers.csv.zip
! unzip /content/sample_submission.csv.zip
! unzip /content/transactions_train.csv.zip
```

```

Archive: /content/articles.csv.zip
  inflating: articles.csv
Archive: /content/customers.csv.zip
  inflating: customers.csv
Archive: /content/sample_submission.csv.zip
  inflating: sample_submission.csv
Archive: /content/transactions_train.csv.zip
  inflating: transactions_train.csv

```

```
[ ]: articles = pd.read_csv('/content/articles.csv')
```

```
[ ]: articles.head(3)
```

```
[ ]:
  article_id  product_code      prod_name  product_type_no  product_type_name \
0   108775015      108775      Strap top             253      Vest top
1   108775044      108775      Strap top             253      Vest top
2   108775051      108775  Strap top (1)             253      Vest top

  product_group_name  graphical_appearance_no  graphical_appearance_name \
0  Garment Upper body             1010016      Solid
1  Garment Upper body             1010016      Solid
2  Garment Upper body             1010017      Stripe

  colour_group_code  colour_group_name  perceived_colour_value_id \
0                9      Black                4
1               10      White                3
2               11  Off White                1

  perceived_colour_value_name  perceived_colour_master_id \
0                Dark                5
1                Light                9
2          Dusty Light                9
```

	perceived_colour_master_name	department_no	department_name	index_code	\
0	Black	1676	Jersey Basic	A	
1	White	1676	Jersey Basic	A	
2	White	1676	Jersey Basic	A	

	index_name	index_group_no	index_group_name	section_no	\
0	Ladieswear	1	Ladieswear	16	
1	Ladieswear	1	Ladieswear	16	
2	Ladieswear	1	Ladieswear	16	

	section_name	garment_group_no	garment_group_name	\
0	Womens Everyday Basics	1002	Jersey Basic	
1	Womens Everyday Basics	1002	Jersey Basic	
2	Womens Everyday Basics	1002	Jersey Basic	

	detail_desc
0	Jersey top with narrow shoulder straps.
1	Jersey top with narrow shoulder straps.
2	Jersey top with narrow shoulder straps.

Warning: Total number of columns (25) exceeds max_columns (20) limiting to first (20) columns.

```
[ ]: customers = pd.read_csv('/content/customers.csv')
```

```
[ ]: customers.head(3)
```

```
[ ]:
      customer_id  FN  Active \
0  0000dbacae5abe5e23885899a1fa44253a17956c6d1c3...  NaN    NaN
1  0000423b00ade91418cceaf3b26c6af3dd342b51fd051e...  NaN    NaN
2  000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...  NaN    NaN
```

	club_member_status	fashion_news_frequency	age	\
0	ACTIVE	NONE	49.0	
1	ACTIVE	NONE	25.0	
2	ACTIVE	NONE	24.0	

	postal_code
0	52043ee2162cf5aa7ee79974281641c6f11a68d276429a...
1	2973abc54daa8a5f8ccfe9362140c63247c5eee03f1d93...
2	64f17e6a330a85798e4998f62d0930d14db8db1c054af6...

```
[ ]: transactions = pd.read_csv('/content/transactions_train.csv')
```

```
[ ]: transactions.head(3)
```

```
[ ]:
      t_dat      customer_id  article_id \
0  2018-09-20  000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...  663713001
```

```

1  2018-09-20  000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...  541518023
2  2018-09-20  00007d2de826758b65a93dd24ce629ed66842531df6699...  505221004

```

```

      price  sales_channel_id
0  0.050831                2
1  0.030492                2
2  0.015237                2

```

1.5 Data Analysis and Visualization

1.5.1 Columns and Shape of the dataset

- The articles file contains 105,542 data points with 25 features.
- The customers file contains 1,371,980 data points with 17 features.
- The transactions file contains 1,371,980 data points with 6 features.

```
[ ]: print(f"Number of articles are {articles.shape[0]}")
      print(f"Number of customers are {customers.shape[0]}")
      print(f"Number of transaction are {transactions.shape[0]}")
```

```

Number of articles are 105542
Number of customers are 1371980
Number of transaction are 31788324

```

```
[ ]: %ls
```

```

articles.csv      customers.csv.zip      transactions_train.csv
articles.csv.zip  sample_submission.csv  transactions_train.csv.zip
customers.csv     sample_submission.csv.zip

```

```
[ ]: ## Features in the articles data
      articles.columns
```

```
[ ]: Index(['article_id', 'product_code', 'prod_name', 'product_type_no',
            'product_type_name', 'product_group_name', 'graphical_appearance_no',
            'graphical_appearance_name', 'colour_group_code', 'colour_group_name',
            'perceived_colour_value_id', 'perceived_colour_value_name',
            'perceived_colour_master_id', 'perceived_colour_master_name',
            'department_no', 'department_name', 'index_code', 'index_name',
            'index_group_no', 'index_group_name', 'section_no', 'section_name',
            'garment_group_no', 'garment_group_name', 'detail_desc'],
           dtype='object')
```

```
[ ]: ## Features in the customers data
      customers.columns
```

```
[ ]: Index(['customer_id', 'FN', 'Active', 'club_member_status',
            'fashion_news_frequency', 'age', 'postal_code'],
```

```
dtype='object')
```

```
[ ]: ## Features in the transaction data  
transactions.columns
```

```
[ ]: Index(['t_dat', 'customer_id', 'article_id', 'price', 'sales_channel_id',  
         'year', 'month', 'day', 'product_type_name', 'age', 'age_bucket',  
         'Season'],  
         dtype='object')
```

1.5.2 Identifying missing values

It's important to identify the missing values in the dataset, so that appropriate missing value imputation can be applied. The following function identifies the total number of missing values in all the columns of the dataset.?

```
[ ]: # Function to identify the missing values  
  
def find_missing_features(df):  
    return dict([(feature, df[feature].isnull().sum()) for feature in df.columns_  
    → if df[feature].isnull().sum() >=1])
```

```
[ ]: find_missing_features(articles)
```

```
[ ]: {'detail_desc': 416}
```

```
[ ]: find_missing_features(customers)
```

```
[ ]: {'Active': 907576,  
     'FN': 895050,  
     'age': 15861,  
     'club_member_status': 6062,  
     'fashion_news_frequency': 16009}
```

```
[ ]: find_missing_features(transactions)
```

```
[ ]: {}
```

- The 'detail_desc' column in the articles dataset contain 416 missing values.
- There are some missing values in the 'Active', 'FN', 'age', 'club_member_status', 'fashion_news_frequency' columns in the customers dataset contain 416 missing records.
- There are no missing values in the transaction table.

1.5.3 Outlier Detection

For outlier detection, we used box plot to visually explore the feature and created a table with different quantile values of the attribute. As we can see in the box plot and the quantile values in the table, there are several outliers in the dataset for 'price' feature which is right skewed in distribution as well.

Identifying the numeric and non-numeric columns

```
[ ]: t_cols = transactions.dtypes[transactions.dtypes.isin([int, float])]
     t_num_cols = list(t_cols.index)
     t_num_cols
```

```
[ ]: ['article_id', 'price', 'sales_channel_id', 'year', 'month', 'day', 'age']
```

```
[ ]: a_cols = articles.dtypes[articles.dtypes.isin([int, float])]
     a_num_cols = list(a_cols.index)
     a_num_cols
```

```
[ ]: ['article_id',
      'product_code',
      'product_type_no',
      'graphical_appearance_no',
      'colour_group_code',
      'perceived_colour_value_id',
      'perceived_colour_master_id',
      'department_no',
      'index_group_no',
      'section_no',
      'garment_group_no']
```


```
[ ]: c_cols = customers.dtypes[customers.dtypes.isin([int, float])]
     c_num_cols = list(c_cols.index)
     c_num_cols
```

```
[ ]: ['FN', 'Active', 'age']
```

Boxplot to visualize outliers

```
[ ]: plt.boxplot(transactions['price'])
     plt.xticks([1], ["Price"])
```

```
[ ]: ([<matplotlib.axis.XTick at 0x7f48b2494710>], [Text(0, 0, 'Price')])
```



Obtaining the percentile values

```
[ ]: def quantile_out(df, cols):
      perc = pd.DataFrame()
      perc = perc.append({'column' : cols, 'min val': min(df[cols]), 'perc 5th' :
      ↪ float(df[cols].quantile(0.05)), 'perc 95th' : float(df[cols].quantile(0.
      ↪ 95)), 'max val': max(df[cols])},
      ignore_index = True)
      return perc
```

```
[ ]: quantile_out(transactions, 'price')
```

```
[ ]:   column  min val  perc 5th  perc 95th  max val
0  price  0.000017   0.00761   0.059305  0.591525
```

Observation: The 95th percentile value of 'Price' is 0.059, whereas the max value is 0.59. This indicates presence of some outliers, or high price items.

```
[ ]: quantile_out(transactions, 'age')
```

```
[ ]:   column  min val  perc 5th  perc 95th  max val
0    age      16.0     21.0     59.0     99.0
```


Observation: There are no outliers in the 'age' variable.

Histogram of age variable

```
[ ]: plt.hist(customers['age'])
```



```
[ ]: (array([3.57169e+05, 3.46715e+05, 1.46283e+05, 1.87960e+05, 1.96469e+05,
            8.39560e+04, 3.15830e+04, 5.38800e+03, 5.19000e+02, 7.70000e+01]),
      array([16. , 24.3, 32.6, 40.9, 49.2, 57.5, 65.8, 74.1, 82.4, 90.7, 99. ]),
      <a list of 10 Patch objects>)
```



1.5.4 Exploratory Data Analysis of Columns

Most of the columns in our dataset are categorical variables. Therefore, we will need to identify the cardinality and the frequency of each categories in the column.

The following function creates a table, barplot, and pie chart to infer the cardinality and frequency.

```
[ ]: def plot_data(data, column):
      """
      Function that takes in dataset and column as input and plots
      frequency bar chart and pie chart for the same
      """
      counts = dict(data[column].value_counts())
      print(data[column].value_counts())
      print()

      figure = plt.figure(figsize=(11,5))

      # bar chart
      plt.subplot(1,2,1)
      ax = sns.countplot(x = data[column])
```

```
ax.set_xticklabels(ax.get_xticklabels(),rotation = 90)

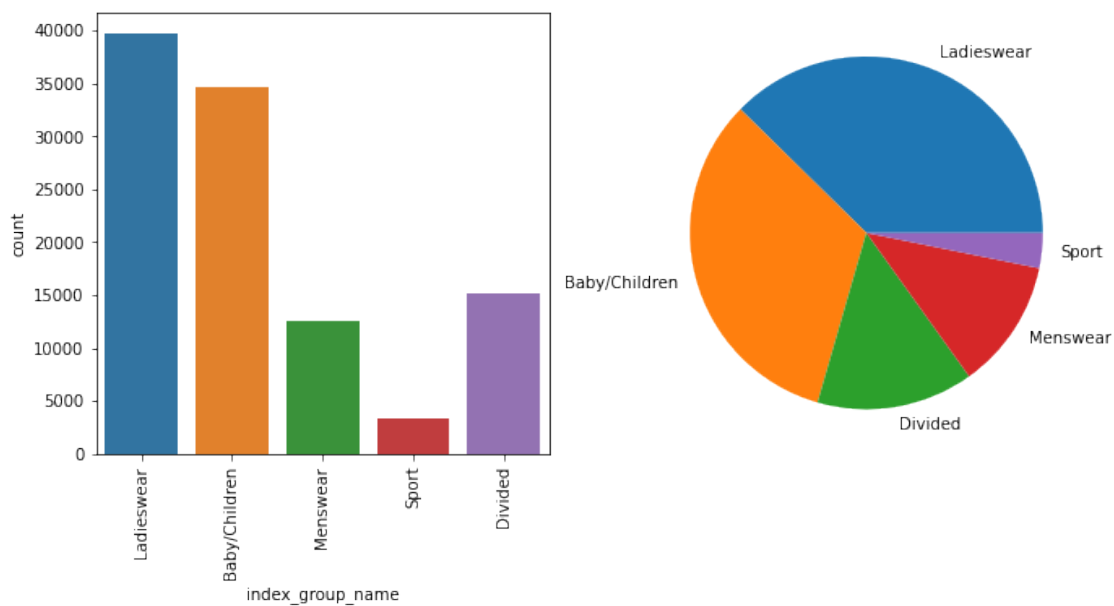
# pie chart
plt.subplot(1,2,2)
plt.pie(x=counts.values(), labels=counts.keys())

plt.show()
```

Exploring different columns of articles

```
[ ]: plot_data(data=articles, column='index_group_name')
```

```
Ladieswear      39737
Baby/Children   34711
Divided         15149
Menswear        12553
Sport           3392
Name: index_group_name, dtype: int64
```




Observation: We see that the ‘Ladieswear’ and ‘Baby/Children’ dominate in the index groups.

```
[ ]: plot_data(data=articles, column='garment_group_name')
```

```
Jersey Fancy      21445
Accessories       11519
Jersey Basic       8126
Knitwear           7490
Under-, Nightwear  7441
```

Trousers	6727
Blouses	5838
Shoes	5145
Dresses Ladies	4874
Outdoor	4501
Unknown	3873
Trousers Denim	3100
Swimwear	2787
Socks and Tights	2272
Shirts	2116
Woven/Jersey/Knitted mix Baby	1965
Shorts	1559
Dresses/Skirts girls	1541
Skirts	1254
Special Offers	1061
Dressed	908

Name: garment_group_name, dtype: int64





Observation: We see that maximum number of article belong to ‘Jersey Fancy’ garment group.

```
[ ]: plot_data(data=articles, column='perceived_colour_master_name')
```

Black	22585
Blue	18469
White	12665
Pink	9403

Grey	8924
Red	5878
Beige	5657
Green	3526
Khaki green	3181
Yellow	3121
Orange	2734
Brown	2269
Metal	2180
Turquoise	1829
Mole	1223
Lilac Purple	1100
Unknown	685
undefined	105
Yellowish Green	5
Bluish Green	3

Name: perceived_colour_master_name, dtype: int64



Observation: The majority of articles are in 'White', 'Blue' or 'Black' colour.

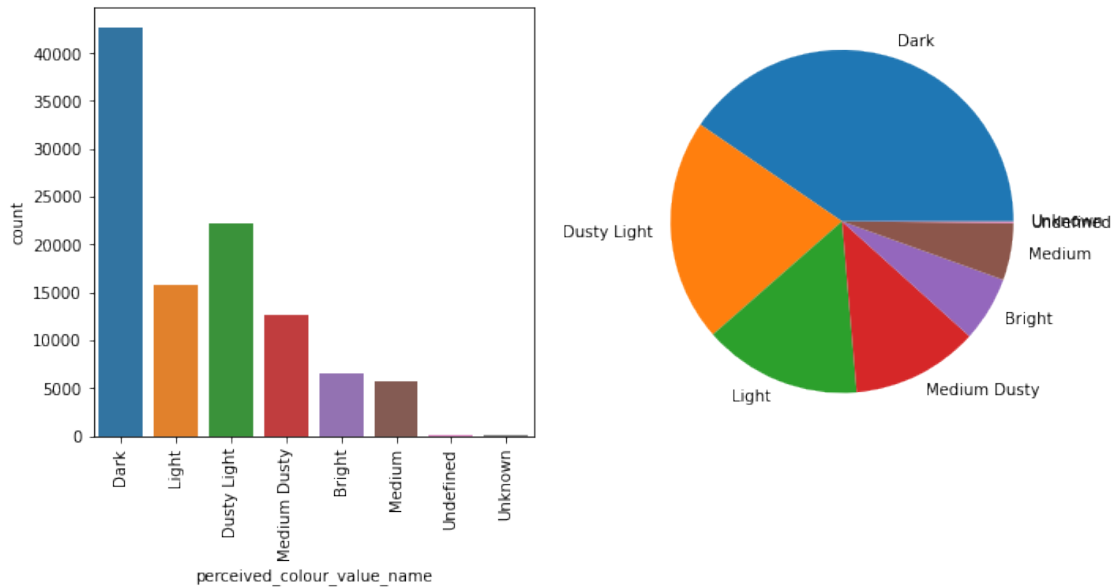
```
[ ]: plot_data(data=articles, column='perceived_colour_value_name')
```

Dark	42706
Dusty Light	22152
Light	15739
Medium Dusty	12630
Bright	6471
Medium	5711

```

Undefined      105
Unknown        28
Name: perceived_colour_value_name, dtype: int64

```



Observation: Majority of articles are 'Dusty Light' or 'Dark' in shade.

```
[ ]: plot_data(data=articles, column='graphical_appearance_name')
```


```

Solid          49747
All over pattern 17165
Melange        5938
Stripe        4990
Denim          4842
Front print    3215
Placement print 3098
Check          2178
Colour blocking 1830
Lace           1513
Other structure 1502
Application/3D 1341
Embroidery     1165
Mixed solid/pattern 1132
Glittering/Metallic 958
Jacquard       830
Sequin         806
Dot            681
Treatment      586
Other pattern  515

```

Contrast	376
Metallic	346
Chambray	322
Slub	153
Transparent	86
Mesh	86
Neps	66
Unknown	52
Argyle	15
Hologram	8

Name: graphical_appearance_name, dtype: int64




Observation: Majority of the articles have a 'Solid' graphical apperance.

```
[ ]: plot_data(data=articles, column='product_group_name')
```

Garment Upper body	42741
Garment Lower body	19812
Garment Full body	13292
Accessories	11158
Underwear	5490
Shoes	5283
Swimwear	3127
Socks & Tights	2442
Nightwear	1899
Unknown	121
Underwear/nightwear	54

Cosmetic	49
Bags	25
Items	17
Furniture	13
Garment and Shoe care	9
Stationery	5
Interior textile	3
Fun	2


Name: product_group_name, dtype: int64



Observation: Most of the articles are of ‘Garment’ product group.

Exploring different columns of customer

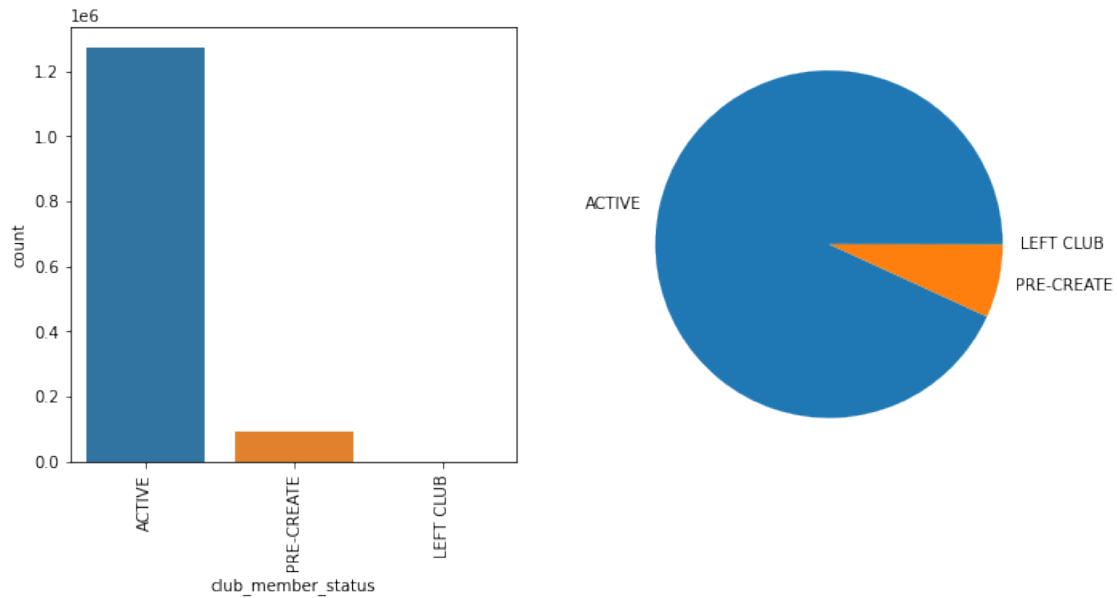
```
[ ]: sns.kdeplot(customers['age'], shade=True)
plt.title('Customer Age Distribution')
plt.show()
```



Observation: The Age of the customers is right skewed and the majority of customers are of the age 18-30 years old.

```
[ ]: plot_data(data=customers, column='club_member_status')
```

```
ACTIVE      1272491
PRE-CREATE   92960
LEFT CLUB     467
Name: club_member_status, dtype: int64
```





Observation: The majority of customers are an 'ACTIVE' club member

Exploring different columns of transaction

```
[ ]: plot_data(data=transactions, column='sales_channel_id')
```

```
2    22379862
1     9408462
Name: sales_channel_id, dtype: int64
```



Observation: Majority of sales are from channel 2.

1.6 Feature Engineering

```
[ ]: transactions["t_dat"] = pd.to_datetime(transactions["t_dat"])
```

```
# add year, month, day
transactions["year"] = transactions["t_dat"].dt.year
transactions["month"] = transactions["t_dat"].dt.month
transactions["day"] = transactions["t_dat"].dt.day
```

```
[ ]: transactions = pd.merge(transactions, articles[["article_id", "product_type_name"]],
                             ↪on="article_id")
```

```
[ ]: transactions = pd.merge(transactions, customers[["customer_id", "age"]],
                             ↪on="customer_id")
display(transactions.head())
```

	t_dat	customer_id	article_id	\
0	2018-09-20	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	663713001	
1	2018-09-24	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	663713001	
2	2018-09-20	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	541518023	
3	2019-03-01	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	578020002	
4	2020-02-03	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	351484002	

	price	sales_channel_id	year	month	day	product_type_name	age
0	0.050831	2	2018	9	20	Underwear body	24.0
1	0.050831	2	2018	9	24	Underwear body	24.0
2	0.030492	2	2018	9	20	Bra	24.0
3	0.013542	2	2019	3	1	Blouse	24.0
4	0.022017	2	2020	2	3	Swimwear bottom	24.0

```
[ ]: bins = [i for i in range(10, 101, 10)]
labels = [i for i in range(1, len(bins))]

transactions["age_bucket"] = pd.cut(transactions["age"], bins=bins,
↪labels=labels)
```

```
[ ]: customer_prod_count = transactions.groupby(["age_bucket"])["product_type_name"].
↪value_counts()
```

```
[ ]: customer_prod = pd.DataFrame(
    index=np.sort(np.array(transactions["age_bucket"].unique().dropna())),
    columns=articles["product_type_name"].unique()
)

customer_prod = customer_prod.fillna(0)
```

```
# count product data
for age_bucket in customer_prod_count.index.get_level_values("age_bucket").
    ↳unique():
    for prod in customer_prod_count.loc[age_bucket].index.
        ↳get_level_values("product_type_name"):
            customer_prod.loc[age_bucket, prod] = customer_prod_count.
                ↳loc[age_bucket].loc[prod]
```

```
[ ]: customer_prod
```

```
[ ]: Vest top    Bra    Underwear Tights    Socks    Leggings/Tights    Sweater \
1      69450    82721                5316    21906                24994    124372
2      621400    644575                90789    196438                306898    1127013
3      280552    240054                48179    85684                 157863    455679
4      219879    185956                28739    88739                 125154    472188
5      177485    152281                22331    74527                 94291    452598
6      34406     21292                4648     11809                 18344    113950
7      4357      2194                585      1349                 2730     22522
8      343       191                30       110                  210      1612
9      49        43                9        43                   23       131

      Top    Trousers    Hair clip    Umbrella    Pyjama jumpsuit/playsuit    Bodysuit \
1      71595    167941        1924        205                210        5510
2      657560    1620163        17561        2158                4326        71332
3      283205    740297         8151         759                4208        33146
4      247712    762682         6052         773                2226        13621
5      244616    704037         5120         873                1157        11345
6      60289     171140         918         220                407         2471
7      10683     30180         137         43                 51         325
8      634      1818          7          2                  3          17
9      77       172          0          0                  0          3

      Hair string    Unknown    Hoodie    ...    Cross-body bag    Moccasins    Towel \
1          2613      4780     37037    ...          110          0        1
2          19958     41032    185979    ...          406         15        7
3           7678     18603     68591    ...          112         16       10
4           7576     16496    101364    ...          141          6        1
5           5185     13431     74196    ...          109          2        1
6           754      2061     11588    ...           21          2        0
7           121       171     1739    ...           6           0        0
8            9        12      178    ...           0           0        0
9            2         3       30    ...           0           0        0

      Wood balls    Zipper head    Mobile case    Pre-walkers    Toy    Marker pen    Bumbag \
1            0          0          236          0      0          35        1
2            5         12          918          0      2          91        8
```

3	10	16	201	0	1	50	2
4	11	13	273	0	2	21	4
5	8	17	184	0	0	28	0
6	1	6	15	0	0	0	0
7	0	0	3	0	0	0	1
8	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0

	Dog wear	Eyeglasses	Wireless earphone case	Stain remover spray	\
1	5	4	62	0	
2	105	9	116	18	
3	64	5	19	5	
4	39	5	45	5	
5	42	1	22	5	
6	9	1	0	3	
7	0	0	0	0	
8	0	0	0	0	
9	0	0	0	0	

	Clothing mist
1	0
2	0
3	1
4	3
5	0
6	0
7	0
8	0
9	0

[9 rows x 131 columns]

```
[ ]: f, ax = plt.subplots(nrows=9, ncols=1, figsize=(12, 24))
ax = ax.flatten()

top_products = []










for i, bin_age in enumerate(customer_prod.index):
    tmp_df = customer_prod.loc[bin_age]

    # descending sort
    indices = tmp_df.values.argsort()[::-1]

    # extract top 10
    columns = customer_prod.columns[indices][:10]
    top_products += columns.tolist()
    ax[i].bar(columns, tmp_df[columns])
```

```
ax[i].set_xticklabels(columns)
ax[i].set_title(f"Top 10 Product Types for the Age Group of_
↳{int(bin_age)*10}'s")

top_products = set(top_products)
plt.tight_layout()
plt.show()
```



```
[ ]: transactions['Season'] = np.where(transactions['month'].isin([3,4,5]),
    ↪ "Summer", "Others")
```

1.7 Encoding and Scaling

Exploring the ‘Transactions’ table

```
[ ]: transactions.info()
transactions['sales_channel_id'].value_counts()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31788324 entries, 0 to 31788323
Data columns (total 5 columns):
#   Column          Dtype
---  -
0   t_dat           object
1   customer_id     object
2   article_id      int64
3   price           float64
4   sales_channel_id int64
dtypes: float64(1), int64(2), object(2)
memory usage: 1.2+ GB
```

```
[ ]: 2    22379862
     1    9408462
     Name: sales_channel_id, dtype: int64
```

Exploring ‘Customers’ To explore the distribution of the features grouped together in order to grasp any tangible information about these features.

```
[ ]: customers[['Active', 'club_member_status', 'fashion_news_frequency']]
    ↪ value_counts()
```

```
[ ]: Active  club_member_status  fashion_news_frequency
1.0      ACTIVE                Regularly                457229
        PRE-CREATE            Regularly                5567
        ACTIVE                Monthly                  735
        PRE-CREATE            NONE                    488
        PRE-CREATE            Monthly                  58
        LEFT CLUB              NONE                     6
        LEFT CLUB              Regularly                 3
dtype: int64
```

1.7.1 Replacing ‘None’ with ‘NONE’ in the *fashion_news_frequency* attribute.

We found a data point with value ‘None’ in the ‘*fashion_news_frequency*’ feature. To congregate this data point with other ‘NONE’ values of the same feature, we are replacing it with ‘NONE’.

```
[ ]: customers.loc[customers['fashion_news_frequency'] == 'None',\
    ↳'fashion_news_frequency'] = 'NONE'
```

To explore the distribution of the features grouped together in order to grasp any tangible information about these features.

```
[ ]: customers[['club_member_status', 'fashion_news_frequency']].value_counts()
```

```
[ ]: club_member_status  fashion_news_frequency
ACTIVE                  NONE                    788484
                        Regularly                471304
PRE-CREATE              NONE                    85065
                        Regularly                5787
ACTIVE                  Monthly                 778
LEFT CLUB               NONE                    459
PRE-CREATE              Monthly                 59
LEFT CLUB               Regularly                 8
dtype: int64
```

1.7.2 One Hot Encoding

Encoding ‘Index Code’ in Articles table There are no ordinal variables in the data set since all the categorical variables in this file do not have any ranking/ordering in between themselves respectively.

Therefore, we are going to employ One Hot Encoding to encode the ‘Index Code’ feature containing around 10 index codes.

```
[ ]: # Defining one-hot encoder object
encoder = OneHotEncoder(sparse = True)

# Performing the encoding
articles_new = pd.DataFrame(encoder.fit_transform(articles[['index_code']]).
    ↳toarray())
articles_new.columns = encoder.get_feature_names_out(['index_code'])

articles = articles.join(articles_new)
```

To check the successful encoding action on the feature in the data set.

```
[ ]: articles.head(3)
```

```
[ ]:   article_id  product_code  prod_name  product_type_no  product_type_name \
0   108775015      108775    Strap top             253      Vest top
1   108775044      108775    Strap top             253      Vest top
2   108775051      108775  Strap top (1)             253      Vest top

   product_group_name  graphical_appearance_no  graphical_appearance_name \
0  Garment Upper body              1010016              Solid
```


1	Garment Upper body	1010016	Solid
2	Garment Upper body	1010017	Stripe

	colour_group_code	colour_group_name	perceived_colour_value_id	\
0	9	Black	4	
1	10	White	3	
2	11	Off White	1	

	perceived_colour_value_name	perceived_colour_master_id	\
0	Dark	5	
1	Light	9	
2	Dusty Light	9	

	perceived_colour_master_name	department_no	...	section_no	\
0	Black	1676	...	16	
1	White	1676	...	16	
2	White	1676	...	16	

	section_name	garment_group_no	garment_group_name	\
0	Womens Everyday Basics	1002	Jersey Basic	
1	Womens Everyday Basics	1002	Jersey Basic	
2	Womens Everyday Basics	1002	Jersey Basic	

	detail_desc	index_code_A	index_code_B	\
0	Jersey top with narrow shoulder straps.	1.0	0.0	
1	Jersey top with narrow shoulder straps.	1.0	0.0	
2	Jersey top with narrow shoulder straps.	1.0	0.0	

	index_code_C	index_code_D	index_code_F	index_code_G	index_code_H	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	

	index_code_I	index_code_J	index_code_S
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0

[3 rows x 35 columns]

Encoding ‘Club Member Status’ and ‘Fashion News Frequency’ variables in Customers table For the features mentioned above, we are going to employ One Hot Encoding to encode their nominal categorical features.

```
[ ]: # Defining one-hot encoder object
encoder = OneHotEncoder(sparse = True)
```

```
# Performing the encoding
customers_new = pd.DataFrame(encoder.
    ↳fit_transform(customers[['club_member_status', 'fashion_news_frequency']]).
    ↳toarray())
customers_new.columns = encoder.get_feature_names_out(['club_member_status', 'fashion_news_frequency'])

customers = customers.join(customers_new)
```

To check the successful encoding action on the feature in the data set.

```
[ ]: customers.head(3)
```

```
[ ]:
      customer_id  FN  Active \
0  0000dbacae5abe5e23885899a1fa44253a17956c6d1c3...  NaN    NaN
1  0000423b00ade91418cceaf3b26c6af3dd342b51fd051e...  NaN    NaN
2  000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...  NaN    NaN

      club_member_status  fashion_news_frequency  age \
0             ACTIVE             NONE  49.0
1             ACTIVE             NONE  25.0
2             ACTIVE             NONE  24.0

      postal_code \
0  52043ee2162cf5aa7ee79974281641c6f11a68d276429a...
1  2973abc54daa8a5f8ccfe9362140c63247c5eee03f1d93...
2  64f17e6a330a85798e4998f62d0930d14db8db1c054af6...

      club_member_status_ACTIVE  club_member_status_LEFT  CLUB \
0             1.0             0.0
1             1.0             0.0
2             1.0             0.0

      club_member_status_PRE-CREATE  club_member_status_nan \
0             0.0             0.0
1             0.0             0.0
2             0.0             0.0

      fashion_news_frequency_Monthly  fashion_news_frequency_NONE \
0             0.0             1.0
1             0.0             1.0
2             0.0             1.0

      fashion_news_frequency_Regularly  fashion_news_frequency_nan
0             0.0             0.0
1             0.0             0.0
2             0.0             0.0
```

1.7.3 Scaling - Normalization or Standardization?

Before proceeding further with the scaling process, we wanted to explore the data set for features which require scaling applied to them.

We also wanted to figure out which scaling techniques (Normalization or Standardization) to apply to the features based on their distribution and nature.

```
[ ]: transactions.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31788324 entries, 0 to 31788323
Data columns (total 5 columns):
#   Column          Dtype
---  -
0   t_dat           object
1   customer_id     object
2   article_id      int64
3   price           float64
4   sales_channel_id int64
dtypes: float64(1), int64(2), object(2)
memory usage: 1.2+ GB
```

```
[ ]: transactions.describe()
```

```
[ ]:
      article_id      price  sales_channel_id
count  3.178832e+07  3.178832e+07  3.178832e+07
mean    6.962272e+08  2.782927e-02  1.704028e+00
std     1.334480e+08  1.918113e-02  4.564786e-01
min     1.087750e+08  1.694915e-05  1.000000e+00
25%     6.328030e+08  1.581356e-02  1.000000e+00
50%     7.145820e+08  2.540678e-02  2.000000e+00
75%     7.865240e+08  3.388136e-02  2.000000e+00
max     9.562170e+08  5.915254e-01  2.000000e+00
```

To check if '*Price*' feature is normally distributed or not

```
[ ]: stats.probplot(transactions['price'], dist="norm", plot=pylab)
pylab.show()
```



Since the data points are deviating significantly from the straight red-line, the ‘*price*’ feature does not seem to be normally distributed

1.7.4 Normalization

We’re going to apply Normalization scaling technique to the ‘*price*’ feature of the transactions table in order to normalize it and scale up the data points within a range.

```
[ ]: # Defining MinMax scaler object
      scaler = MinMaxScaler()

      # Performing the scaling
      transactions['price_scaled'] = pd.DataFrame(scaler.
      ↪fit_transform(transactions[['price']]))

[ ]: transactions.head(5)
```

```
[ ]:      t_dat      customer_id  article_id  \
0  2018-09-20  000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...  663713001
1  2018-09-20  000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...  541518023
2  2018-09-20  00007d2de826758b65a93dd24ce629ed66842531df6699...  505221004
3  2018-09-20  00007d2de826758b65a93dd24ce629ed66842531df6699...  685687003
4  2018-09-20  00007d2de826758b65a93dd24ce629ed66842531df6699...  685687004
```

	price	sales_channel_id	price_scaled
0	0.050831	2	0.085905
1	0.030492	2	0.051520
2	0.015237	2	0.025731
3	0.016932	2	0.028597
4	0.016932	2	0.028597

```
[ ]: transactions.describe()
```

```
[ ]:
      article_id      price  sales_channel_id  price_scaled
count  3.178832e+07  3.178832e+07      3.178832e+07  3.178832e+07
mean    6.962272e+08  2.782927e-02      1.704028e+00  4.701932e-02
std     1.334480e+08  1.918113e-02      4.564786e-01  3.242748e-02
min     1.087750e+08  1.694915e-05      1.000000e+00  0.000000e+00
25%     6.328030e+08  1.581356e-02      1.000000e+00  2.670564e-02
50%     7.145820e+08  2.540678e-02      2.000000e+00  4.292387e-02
75%     7.865240e+08  3.388136e-02      2.000000e+00  5.725092e-02
max     9.562170e+08  5.915254e-01      2.000000e+00  1.000000e+00
```

1.8 What intrigued you about the data? Why does that matter?

The data appeared in a kaggle competition to recommend the customers clothing basis their demographics and past purchase data.

Recommendor systems are a norm now in various industries like OTT and FMCG. However, it is a little tricky for the fashion industry to predict customer future purchase as the trend keeps on changing rapidly and unlike the Groceries in FMCG. Customers are less likely to buy the same clothing again. Further, we have the image data of the products and the purchase behaviour is also affected by the images of the products available on the website.

We took this data to get familiar with the data and customer behaviour in the clothing industry.

1.9 What would your proposed next steps be?

Our next steps would be exploring the image data and how it can add value to the recommendor system we will be creating. We have to decide on the what model we will chose for this use case basis this EDA and our further exploration of the image data.

Further, the data size is large (>30 GB). Hence we need to decide on how to get enough processing power to create model with this data.

1.10 Observations and Findings

In our initial analysis, we started off by checking for any missing data. We found several data points in the ‘Articles’ and ‘Customers’ data set missing information which were dealt by imputing them by either their central tendencies or were removed, if insignificant.

We moved on to check the cardinality of the features and found out that some of them have very high cardinality. We performed some necessary actions to deal with or remove them accordingly.

Next, we used techniques to check for outliers in all the 3 data sets and found several outliers present in the ‘*transactions*’ data set. Necessary steps were performed to remove outliers from the data set.

Later, we implemented feature engineering techniques to extract and modify date-time information from ‘*transactions*’ table and performed exploratory data analysis (EDA) to find insights about the pattern and behaviour of customers with their past purchases.